Assignment 1: Evil Hangman

Due Thursday, November 15, 11:59 PM

Introduction

It's hard to write computer programs to play games. When we humans sit down to play games, we can draw on past experience, adapt to our opponents' strategies, and learn from our mistakes. Computers, on the other hand, blindly follow a preset algorithm that (hopefully) causes it to act intelligently. Though computers have bested their human masters in some games, most notably checkers and chess, the programs that do so often draw on hundreds of years of human experience and use extraordinarily complex algorithms and optimizations to outcalculate their opponents.

While there are many viable strategies for building competitive computer game players, there is one approach that has been fairly neglected in modern research – cheating. Why spend all the effort trying to teach a computer the nuances of strategy when you can simply write a program that plays dirty and wins handily all the time? In this assignment, you will build a mischievous program that bends the rules of *Hangman* to trounce its human opponent time and time again. In doing so, you'll cement your skills with the STL containers and will hone your general programming savvy. Plus, you'll end up with a highly entertaining piece of software, at least from your perspective. ☺

In case you aren't familiar with the game *Hangman*, the rules are as follows:

1. One player chooses a secret word, then writes out a number of dashes equal to the word length.

2. The other player begins guessing letters. Whenever she guesses a letter contained in the hidden word, the first player reveals each instance of that letter in the word. Otherwise, the guess is wrong.

3. The game ends either when all the letters in the word have been revealed or when the guesser has run out of guesses.

Fundamental to the game is the fact the first player accurately represents the word she has chosen. That way, when the other players guess letters, she can reveal whether that letter is in the word. But what happens if the player doesn't do this? This gives the player who chooses the hidden word an enormous advantage. For example, suppose that you're the player trying to guess the word, and at some point you end up revealing letters until you arrive at this point with only one guess remaining:

D O – B L E

There are only two words in the English language that match this pattern: "doable" and "double." If the player who chose the hidden word is playing fairly, then you have a fifty-fifty chance of winning this game if you guess 'A' or 'U' as the missing letter. However, if your opponent is cheating and hasn't actually committed to either word, then there is no possible way you can win this game. No matter what letter you guess, your opponent can claim that she had picked the other word, say that your guess is incorrect, and win the game. That is, if you guess that the word is "doable," she can pretend that she committed to "double" the whole time, and vice-versa.

Let's illustrate this technique with an example. Suppose that you are playing *Hangman* and it's your turn to choose a word, which we'll assume is of length four. Rather than committing to a secret word, you instead compile a list of every four-letter word in the English language. For simplicity, let's assume that English only has a few four-letter words, all of which are reprinted here:

<p align="center">ALLY BETA COOL DEAL ELSE FLEW GOOD HOPE IBEX</p>

Now, suppose that your opponent guesses the letter 'E.' You now need to tell your opponent which letters in the word you've "picked" are E's. Of course, you haven't picked a word, and so you have multiple options about where you reveal the E's. Here's the above word list, with E's highlighted in each word:

<p align="center">ALLY B**E**TA COOL D**E**AL **E**LS**E** FL**E**W GOOD HOP**E** IB**E**X</p>

If you'll notice, every word in your word list falls into one of five "word families:"

- ----, which contains the word ALLY, COOL, and GOOD.
- -E--, containing BETA and DEAL.
- --E-, containing FLEW and IBEX.
- E--E, containing ELSE.
- ---E, containing HOPE.

Since the letters you reveal have to correspond to *some* word in your word list, you can choose to reveal any one of the above five families. There are many ways to pick which family to reveal – perhaps you want to steer your opponent toward a smaller family with more obscure words, or toward a larger family in the hopes of keeping your options open. For this assignment, in the interests of simplicity, we'll adopt the latter approach and always choose the largest of the remaining word families. In this case, it means that you should pick the family ----. This reduces your word list down to

<p align="center">ALLY COOL GOOD</p>

and since you didn't reveal any letters, you would tell your opponent that his guess was wrong.

Let's see a few more examples of this strategy. Given this three-word word list, if your opponent guesses the letter O, then you would break your word list down into two families:

- -OO-, containing `COOL` and `GOOD`.
- ----, containing `ALLY`.

The first of these families is larger than the second, and so you choose it, revealing two O's in the word and reducing your list down to

COOL   GOOD

But what happens if your opponent guesses a letter that doesn't appear anywhere in your word list? For example, what happens if your opponent now guesses 'T'? This isn't a problem. If you try splitting these words apart into word families, you'll find that there's only one family: the family ---- containing both `COOL` and `GOOD`. Since there is only one word family, it's trivially the largest, and by picking it you'd maintain the word list you already had.

There are two possible outcomes of this game. First, your opponent might be smart enough to pare the word list down to one word and then guess what that word is. In this case, you should congratulate him – that's an impressive feat considering the scheming you were up to! Second, and by far the most common case, your opponent will be completely stumped and will run out of guesses. When this happens, you can pick any word you'd like from your list and say it's the word that you had chosen all along. The beauty of this setup is that your opponent will have no way of knowing that you were dodging guesses the whole time – it looks like you simply picked an unusual word and stuck with it the whole way.

The Assignment

Using only standard C++ (that is, **without using any of the CS106B/X libraries**), your assignment is to write a computer program which plays a game of *Hangman* using this "Evil Hangman" algorithm. In particular, your program should do the following:

1. Read the file `dictionary.txt`, which contains the master word list.

2. Prompt the user for a word length, prompting as necessary until she enters a number such that there's at least one word that's exactly that long. That is, if the user wants to play with words of length -42 or 137, since no English words are that long, you should reprompt her.

3. Prompt the user for a number of guesses, which must be an integer greater than zero. Don't worry about unusually large numbers of guesses – after all, having more than 26 guesses is not going to help your opponent!

4. Prompt the user for whether she wants to have a running total of the number of words remaining in the word list. This completely ruins the illusion of a fair game that you'll be cultivating, but it's quite useful for testing (and grading!)

5. Play a game of *Hangman* using the Evil Hangman algorithm, as described below:

1. Construct a list of all words in the English language whose length matches the input length.

2. Print out how many guesses the user has remaining, along with any letters the player has guessed and the current blanked-out version of the word. If the user chose earlier to see the number of words remaining, print that out too.

3. Prompt the user for a single letter guess, reprompting until the user enters a letter that she hasn't guessed yet. Make sure that the input is exactly one character long and that it's a letter of the alphabet.

4. Partition the words in the dictionary into groups by word family.

5. Find the most common "word family" in the remaining words, remove all words from the word list that aren't in that family, and report the position of the letters (if any) to the user. If the word family doesn't contain any copies of the letter, subtract a remaining guess from the user.

6. If the player has run out of guesses, pick a word from the word list and display it as the word that the computer initially "chose."

7. If the player correctly guesses the word, congratulate her.

6. Ask if the user wants to play again and loop accordingly.

It's up to you to think about how you want to partition words into word families. Think about what data structures would be best for tracking word families and the master word list. Would a `vector` work? How about a `map`? A `stack` or `queue`? Thinking through the design before you start coding will save you a lot of time and headache.

Efficiency Concerns

The program you will be writing maintains a grand illusion: it pretends to play a game of *Hangman*, but instead does something more nefarious behind-the-scenes. Consequently, you should try to make your program as responsive as possible. If players have to wait three or more seconds after entering a letter, they're almost certain to notice that something is awry. The illusion will be broken, and the beauty of your program will be lost.

Although efficiency is important in *Evil Hangman*, never forget the wisdom of programmers past:[*]

Premature optimization is the root of all evil.

There is an inherent tradeoff between *readable code* and *optimized code*. Code that is clean and easy to follow often contains subtle sources of inefficiency, while

---

[*] This has been attributed to both Don Knuth and Tony Hoare, both of whom are legends in the field, so it should carry some weight!

code that is optimized is often impossibly difficult to read. As an extreme example, consider the following piece of C code, known as *Duff's Device*:

```
int n = (count + 7) / 8;

switch (count % 8) {
    case 0: do {  *to = *from++;
    case 7:       *to = *from++;
    case 6:       *to = *from++;
    case 5:       *to = *from++;
    case 4:       *to = *from++;
    case 3:       *to = *from++;
    case 2:       *to = *from++;
    case 1:       *to = *from++;
            } while (--n > 0);
}
```

This code[†] is entirely equivalent to this (much more readable) C++ statement:

```
copy(from, from + n, to);
```

On some machines, Duff's Device slightly performs outperforms a hand-written `for` loop or a call to the `copy` algorithm due to the way that the processor handles loops. Of course, Duff's Device is impossibly hard to read, and even the best of program-mers would have no way of deciphering it without extensive commenting.

As you write your Evil Hangman implementation, try to get the program working **before** you worry about runtime efficiency. If you try to optimize the code as you write it, you are almost certainly going to end up with code that is harder to read and maintain, making debugging a nightmare. Moreover, if you optimize as you go, there's no guarantee that your optimization will have any noticeable effect. If you spend a long time optimizing parts of the code that aren't actually contributing to the overall runtime, you'll have complicated the code with little to show for it.

The moral of the story is this: first get your program working, *then* worry about making it run quickly. Chances are, your program will be fast enough the first time you write it and you won't need to optimize it. If for some reason this isn't the case, feel free to send me an email and I can try to offer some advice.

Advice, Tips, and Tricks

There is no starter code for this program, and you'll be building it from scratch. Consequently, you'll need to do a bit of planning to figure out what the best data structures are for the program. There is no "right way" to go about writing this pro-gram, but some design decisions are much better than others (e.g. you *can* store your word list in a `stack` or `map`, but this is probably not the best option). Here are some general tips and tricks that might be useful:

1. *Letter position matters just as much as letter frequency*. When computing word families, it's not enough to count the number of times a particular letter appears in a word; you also have to consider their positions. For example, "BEER" and "HERE" are in two different families even though they both have

---

† Retrieved from http://foldoc.org/Duff%27s+device

two E's in them. Make sure your representation of word families can encode this distinction.

2. *Watch out for gaps in the dictionary*. When the user specifies a word length, you will need to check that there are indeed words of that length in the dictionary. You might initially assume that if the requested word length is less than the length of the longest word in the dictionary, there must be some word of that length. Unfortunately, the dictionary contains a few "gaps." The longest word in the dictionary has length 29, but there are no words of length 26 or 27. Be sure to take this into account when checking if a word length is valid.

3. *Don't explicitly enumerate word families*. If you are working with a word of length n, then there are 2^n possible word families for each letter. However, most of these families don't actually appear in the English language. For example, no English words contain three consecutive U's, and no word matches the pattern `E-EE-EE--E`. Rather than explicitly generating every word family whenever the user enters a guess, see if you can generate word families only for words that actually appear in the word list.

4. *Be careful when iterating over containers while removing elements*. If you remove an element from a container, you *invalidate* all iterators that point to that element. Invalid iterators may point to garbage data, and their `*` and `++` operators are not guaranteed to work as expected. If you're working with a sequence container (`list`, `vector`, or `deque`), you can avoid issues with invalidated iterators by updating the value of the iterator you pass to erase to the iterator returned by `erase`. For example:

```
while(itr != myContainer.end() {
    if (/* some condition */)
        itr = myContainer.erase(itr);
    else
        ++itr;
}
```

When working with associative containers (`map`, `set`), if you have an iterator `itr` to an element that you want to remove, and you're also using `itr` inside of a loop, you can prevent `itr` from being invalidated by using this technique:

```
while(itr != myContainer.end()) {
    if (/* some condition */) {
        containerType::iterator toRemove = itr;
        ++itr;                          // Advance itr.
        myContainer.erase(toRemove); // Remove element.
    }
    /* Otherwise, we didn't advance itr yet, so do it here. */
    else ++itr;
}
```

Also, do be aware that removing elements at arbitrary positions inside of an STL `vector` or `deque` can be very slow. In fact, if you are planning on removing a large number of elements from a `vector` or `deque`, you are better off cre-

ating a brand-new container holding all of the elements you'd like to retain, then overwriting the original `vector` or `deque` with this new one.

Extensions

The algorithm outlined in this handout is by no means optimal, and there are several cases in which it will make bad decisions. For example, suppose that the human has exactly one guess remaining and that computer has the following word list:

<div align="center">DEAL   TEAR   MONK</div>

If the human guesses the letter 'E' here, the computer will notice that the word family `-E--` has two elements and the word family `----` has just one. Consequently, it will pick the family containing `DEAL` and `TEAR`, revealing an E and giving the human another chance to guess. However, since the human has only one guess left, a much better decision would be to pick the family `----` containing `MONK`, causing the human to immediately lose the game.

More generally, picking the largest of the remaining word families is not necessarily the best option. After you implement this assignment, take some time to think over possible improvements to the algorithm. You might weight the word families using some metric other than size. You might consider having the computer "look ahead" a step or two by considering what actions it might take in the future.[*] If you implement something interesting, feel free to include it with your solution... I'd love to see what you've cooked up!

Deliverables

You'll submit this assignment on paperless just like your last one. If you've improved upon or modified the base algorithm in some way, let me know in the comments.

Good luck!

---

[*]   This idea of looking ahead generalizes to a strategy called *minimax*, which can be shown to play a theoretically perfect game. Researching and implementing a minimax search can make your player substantially more powerful and would be an excellent way to show off your CS wizardry.